

T	S	P
warm	sun	0.4
warm	rain	0.1
cold	sun	0.2
cold	rain	0.3

ALWAYS REMEMBER TO NORMALISE

Inference by enumeration. Space Complexity: $O(d^n)$ to store the joint distribution. Worst case time complexity: $O(d^n)$

2. INDEPENDENCE

Two variables are independent iff for all values in the domain of X and Y. $P(X|Y) = P(X)$ and vversa. Helps reduce complexity more instead of full joint.

X is conditionally independent of Y given Z iff for all values in the domain of X, Y and Z:

$$P(X|Y, Z) = P(X|Z)$$

$$P(X, Y|Z) = P(X|Z)P(Y|Z)$$

CPT: conditional probability table for each node.

2.5 CHAIN RULE

Obtain the joint prob by multiplying the conditional probabilities by chain

Generalization of product rule $P(X_1, X_2, X_3, \dots, X_n)$

$$= P(X_1) * P(X_2 | X_1) P(X_3 | X_2, X_1) * \dots * (X_n | X_{n-1} \dots X_1)$$

Bayes Rule $P(X, Y) = P(X|Y)P(Y) = P(Y|X)P(X)$

Big-O notation tells us how algorithm's run-time behave as input size n asymptotically approaches infinity

BIG O ABSTRACTS out constant factors, only keeps the complexity

And ONLY KEEP MOST DOMINANT TIME FACTOR

Querying in Bayesian Networks using joint by chain rule all the CPTs.

X is independent of all the non-descendants, given all its parents (U_i). It is conditionally independent of Zs given its parents.

Note also that when summed across a row, probability is 1. Observe

$$\sum_M P(M|A=T) = P(M=T|A=T) + P(M=F|A=T) = 1;$$

for given fixed A value

$P(b, e, j) = ?$ Note that A, M are hidden variables


$$= \sum_{A, M} P(b, e, A, j, M)$$

$$= \sum_{A, M} P(b)P(e)P(A|b, e)P(j|A)P(M|A) = P(b)P(e) \sum_{A, M} P(A|b, e)P(j|A)P(M|A)$$

$$= P(b)P(e) \sum_A P(A|b, e)P(j|A) \sum_M P(M|A) = P(b)P(e) \sum_A P(A|b, e)P(j|A)$$

$$= P(b)P(e)[P(a|b, e)p(j|a) + P(-a|b, e)p(j|¬a)]$$

Activate Windows



4. BN DSEP

For Markov Blanket rules, variable is independent of other nodes given its direct Markov Blanket -> Parents + Children + Parent of same children. If two nodes are d-separated, they are independent

PATH CASE CHECKS. Find $P(A, C)$ vs $P(A, C | B)$.

Sum out unseen nodes not wanted in query.

D-separation does not guarantee dependence as they can be numerically independent

Analyze the graph by breaking down all the paths connecting the two variables into the 3 cases we have seen so far. If all possible paths are cut off or inactive, the two nodes are called dseparated or independent

MORALISATION

1. Keep ALL ancestor or direct parents of all variables in concern.

2. For each pair of variables with common child, draw a line between them, between the parents. If there are more than two parents, draw edge between each one of them 3. Delete Z (evidence nodes) and its edges.

Independent is guaranteed if: - X and Y are disconnected completely. No possible path to trace.

5. MLE Finds the set of parameters (Θ) that maximise the likelihood (probability) of seeing the data D . $\text{argmax } L(\Theta)$ wrt $\Theta = \text{argmax } P(D|\Theta)$

Obtain the likelihood using the assumption that observations are independent and identically distributed Prob of R.V seeing the value and the times it happens, all independently and multiplied together.

$$\text{argmax}_{\Theta} \log L(\Theta) + \lambda \left(\sum_{i=1}^m \theta_i - 1 \right) \sum_{i=1}^m \theta_i^{n_i}$$

2. To find the value that maximise $L(\Theta)$, apply log transformation. //+lagrange

3. Use the probability constraint of all probs sum =1 and sub it in.

4. Obtain derivatives of $L(\Theta_i)$ with respect to Θ_i and equate to 0

If many probability parameters Θ_n , use LAGRANGE.

Generically the solution is -> $\Theta_i = n_i / \text{total } n \text{ over all } n_i$.

FOR CPT MLE estimate, include all parameters within a CPT for each given probability. Be careful of whether its full or least params CPT.

need to write out the entire log likelihood function across row and down all columns to get all the parameters. $\sum \log(\Theta(b_i | a_i)) =$

$$l(\theta_B) = n_{BA}(b_1, a_1) \cdot \log \theta_B(b_1|a_1) + n_{BA}(b_2, a_1) \cdot \log \theta_B(b_2|a_1) + n_{BA}(b_1, a_2) \cdot \log \theta_B(b_1|a_2) + n_{BA}(b_2, a_2) \cdot \log \theta_B(b_2|a_2)$$

Laplace Smoothing/Add-one Smoothing. Add $\lambda = 1$ to all the freq counts

Higher values of λ pushes the distribution towards uniform! λ values diminishes with more data and higher counts. Remember that plus 1 affects the DENOMINATOR ALSO

Active Triples		Inactive Triples	
linear			
diverging			
converge			
converge on child			

B	#	P(B)
b_1	4+1	5/7
b_2	1+1	2/7

B	D	#	P(D B)
b_1	d_1	1+1	2/6
b_1	d_2	3+1	4/6
b_2	d_1	1+1	2/3
b_2	d_2	0+1	1/3

6. Markov Models — Value or the state of a random variable (X_t) in each time step only depends on the previous time step.

$p(X_{t+1} | X_t \dots X_2 X_1) = \text{only } P(X_{t+1} | X_t)$ Stationarity assumption: transition probabilities between the states are always the same

Joint distribution for a basic markov chain $Y_1 \rightarrow Y_2 \rightarrow \dots \rightarrow Y_n$

$$P(Y_1, Y_2, \dots, Y_n) = P(Y_1) \circ \prod P(Y_i | Y_{i-1})$$

For markov, need INITIAL STATE PROB, TRANSITION PROBS

Probability distribution at any point itself like $P(X_3)$ can be derived

$$\begin{aligned} \text{Example } \rightarrow P(V_3 = h) &= \sum_{V_2} P(V_3 = h, V_2) \\ &= \sum_{V_2} P(V_3 = h | v_2) \circ P(v_2) \\ &= [P_3(h|h) + P_3(h|l)] \circ (\sum_{V_1} P(v_2 | v_1)) \circ P_1(v_1) \\ &= P_3(h|h) * [\sum_{V_1} P(v_2 = h|v_1) * P(v_1)] \\ &\quad + P_3(h|l) * [\sum_{V_1} P(v_2 = h|v_1) * P(v_1)] \end{aligned}$$

For majority of the markov chains, the stationary distribution is independent of the initial distribution. Satisfies distribution at infinity. Remember sub back prob constraints = 1.

Infinity stationary probs -> NOT depend on initial state probs, Only dep on transit probs

$$P_{\infty}(X) = P_{\infty+1}(X)$$

$$P_{\infty}(X) = \sum_x P(X|x)P_{\infty}(x) \text{ (eq. I)}$$

$$\sum_x P_{\infty}(x) = 1 \text{ (eq. II)}$$

$$P_{\infty+1}(V) = \sum_v P(V|v)P_{\infty}(v) = P_{\infty}(V)$$

CPT assumed to not change

$$P_{\infty}(\text{high}) = P(\text{high}|\text{high})P_{\infty}(\text{high}) + P(\text{high}|\text{low})P_{\infty}(\text{low})$$

$$\rightarrow P_{\infty}(\text{high}) = 0.9 * P_{\infty}(\text{high}) + 0.3 * P_{\infty}(\text{low})$$

$$2. \sum_v P_{\infty}(V) = 1 \Rightarrow P_{\infty}(\text{low}) + P_{\infty}(\text{high}) = 1$$

$$\sum_{t=0}^{\infty} \gamma^t R(s_t)$$

12. REINFORCE LEARNING write out all $T(S, A, S')$ & $R(S, A, S')$

The utility function over a path s_0, s_1, s_2, \dots gives the sum of discounted rewards

Goal to learn Optimal policy ($\pi^*(s)$): a function that specifies the action to take in state s that results in the highest expected utility.

Essential -> argmax choose highest util action pick the action that gives the best util outcome. for EACH ACTION write sum all probs and the outcomes of actions, since action can have more than one outcome + reward

Can be like Find all iterations 1 or more $V_1^*(S=\{0,1,2,3\})$ -> find for each iteration each $V_1^*(S)$ for all states.

For next iterations, the V_2^* will use $V_1^*(s)$ for action that results in state s

$V_i^*(s)$ is the estimate of $V^*(s)$ in iteration i

1. Start with $V_0^*(s) = 0 \quad \forall s \in S$

2. Compute V_{i+1}^* given V_i^* :

$$V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')] \text{ (Bellman equation)}$$

3. Repeat until convergence

\rightarrow i.e., $V_{i+1}^*(s)$ is very close to $V_i^*(s) \quad \forall s$

For true optimal V^* , $V^* = R(s) + \gamma V^*$

Subtracting the above two equations:

$$V_{i+1}^* - V^* = \gamma(V_i^* - V^*) \quad (\gamma < 1)$$

Distance between the estimated and optimal would decrease by a factor γ for each iteration

7. HMM includes evidences into Markov

three types of parameters for HMM. Initial $P(X_1)$. Transition $P(X_i|X_{i-1})$. Emission $P(E_i|X_i)$. Evidence observation independent of other evidences.

(DEPENDS ON STRUCTURE OF HMM VARIATION)

$$P(X_1, \dots, X_n, Y_0, Y_1, \dots, Y_n, Y_{n+1})$$

$$= P(Y_0, Y_1, \dots, Y_n, Y_{n+1}) P(X_1, \dots, X_n | Y_0, Y_1, \dots, Y_n, Y_{n+1})$$

$$= P(Y_0) \prod_{i=1}^{n+1} P(Y_i | Y_{i-1}) \prod_{i=1}^n P(X_i | Y_i)$$

$P(Y_i | Y_{i-1})$: transition distribution

$P(X_i | Y_i)$: output/emission/observation distribution

Based on counting. **a_{ij} is the times a for value i transit to value j for a random variable.** Number of times value i appears **does not require it to start at value i, just seeing how many times IT APPEARS.**

$$a_{i \rightarrow j} = \frac{\text{count}(i \rightarrow j)}{\text{count}(i)}, \text{ s.t. count}(i) \text{ is just total number of times value i appear,}$$

$$b_i(o) = \frac{\text{count}(i \rightarrow o)}{\text{count}(i)}$$

8. VITERBI Algo. for (y_1, y_n) sequences, N hidden states $\rightarrow O(N^n)$

N possible sequences, n sequence lengths

$$\Pi(t, 1) = \max_{i=\{1,2,3\}} \Pi(t-1, i) a_{i,1} b_1(x_t)$$

Recursive Equation!

E.g., For $k=2$ and three states

- $\Pi(2,1) = \max(\Pi(1,1)a_{1,1}b_1(x_2), \Pi(1,2)a_{2,1}b_1(x_2), \Pi(1,3)a_{3,1}b_1(x_2))$
- $= \max_{i=\{1,2,3\}} \Pi(1, i) a_{i,1} b_1(x_2)$
- $\Pi(2,2) = \max(\Pi(1,1)a_{1,2}b_2(x_2), \Pi(1,2)a_{2,2}b_2(x_2), \Pi(1,3)a_{3,2}b_2(x_2))$
- $= \max_{i=\{1,2,3\}} \Pi(1, i) a_{i,2} b_2(x_2)$

Use $\Pi(k, v)$ to represent the max probability over all length k sequences ending in state v

Base case:

- $\Pi(1, v) = a_{0,v} b_v(x_1)$

Recursive step (moving forward)

- for $k \in \{2, \dots, n\}$:
- for $v \in \{1, \dots, N-1\}$:

$$\Pi(k, v) = \max_{u \in \{1, \dots, N-1\}} \{ \Pi(k-1, u) a_{u,v} b_v(x_k) \}$$

Last step: $\max_{v \in \{1, \dots, N-1\}} \{ \Pi(n, v) a_{v, \text{STOP}} \}$

$$= \max_{Y_1, Y_2, \dots, Y_n, Y_{n+1}} P(Y_1) P(X_1 | Y_1) P(Y_2 | Y_1) P(X_2 | Y_2) \dots P(Y_n | Y_{n-1}) P(X_n | Y_n) P(Y_{n+1} | Y_n)$$

$$\Pi(1, Y_1 = v) = P(Y_1 = v | Y_0) P(X_1 | Y_1 = v) \text{ (base case)}$$

$$\Pi(k, Y_k = v) = \max_{Y_{k-1}} \Pi(k-1, Y_{k-1}) P(Y_k = v | Y_{k-1}) P(X_k | Y_k = v)$$

(recurrence equation)

9. Variable elimination for networks, By grouping and transforming, compute lesser so we compute once and reuse

$$P(D) = \sum_C \sum_B \sum_A P(D|C) P(C|B) P(A) P(B|A)$$

$$= \sum_C P(D|C) \sum_B P(C|B) \underbrace{\sum_A P(A) P(B|A)}_{\tau_1(B)}$$

$$\tau_2(C)$$

Each factor has a scope \rightarrow all Random variables included in the factor for CPT. 1. Here you see for $P(C|E, F)$ the factor $\phi_C|EF$ contains scope C,E,F \rightarrow 2. $\Psi(X, Y) = \phi_X * \phi_Y$.

Scope Ψ is union of scope of all multiplied factors.

3. $\tau(X) = \text{sumout } Y \sum \Psi(X, Y)$. Scope $\tau(X)$ is exclude eliminated var. Need not be valid prob distributions

Eliminate D

$$\Phi^1 = \{ \phi_L(L, G), \phi_S(S, I), \phi_I(I), \phi_G(G, D, I), \phi_D(D) \}$$

- Join on D: $\phi_1(G, D, I) = \phi_G(G, D, I) \phi_D(D)$
- Sum out D: $\tau_1(G, I) = \sum_D \phi(G, D, I)$
- Update remaining factors: $\Phi^2 = \{ \phi_L(L, G), \phi_S(S, I), \phi_I(I), \phi_G(G, D, I), \phi_D(D), \tau_1(G, I) \}$

Eg Order of elimination maybe given/selected, compared to find fastest way of eliminating. After elim all the variables, **normalise final table ϕ_{final} .**

avoid exponential complexity or have $m \ll n$? Dep on order of elimination, largest factor during elim? depend on largest table for interim factor found ϕ .

VE complexity = $O(n * N_{\text{max}})$, $N_{\text{max}} = k^m$, m vars with k values

VE better than ENUMERATION of $O(n * k^n)$, since N_{max} lesser usually

Dep on size of largest interim factor

10. Approx inference by SAMPLING.

Sampling from multinom distribution. Draw a U value from dist. Convert U to a variable outcome base on subinterval of prob table

Example: Sampling from $P(A)$

A	P(A)
a_0	0.6
a_1	0.1
a_2	0.3

- $0 \leq u < 0.6, \Rightarrow a_0$
- $0.6 \leq u < 0.7, \Rightarrow a_1$
- $0.7 \leq u < 1, \Rightarrow a_2$

Prior sampling Draw several samples from the Bayes' net by visiting each variable in topology order. Prob of sample is consistent as LLN. Weakness for many variables (eg100s), hard to get enough samples of rare events

Rejection Sampling: Intuition can calculate the conditional probability as: \rightarrow Sample from prior distribution (prior sampling) \rightarrow **Ignore (reject) samples that don't match evidence** \rightarrow Use counting to obtain the probability. Only counting amongst those relevant to evidence of query However it can't

- \rightarrow Rejects many samples if the evidence is unlikely!
- \rightarrow Fraction of samples consistent with evidence decreases as the #evidence variables grows.

OR LIKELIHOOD WEIGHTING \rightarrow Idea: fix evidence variables and sample the rest. **Do this when generation values are provided then you can fix evidences. If samples pregiven only consider those consistent with evidences, recall tutorial qn**

Example: weight is 0.8 for the sample [b, a] just weight it by 0.8 rather than explicitly sampling.

Multiply weights only by the evidences var values. LLweight also consistent sampling LLN

Ex: $P(C|S=s, W=w)$

Sampling Process (topological order).

- Sample from $P(C)$, value = c
- Sprinkler is an evidence variable: Value = s, weight = $1 * 0.1$

set initial weight = 1

- Sample from $P(R|C=c)$, value = r
- Wet grass is an evidence variable: Value = w, weight = $0.1 * 0.99$
- Sample returned: [c, s, r, w] with weight = 0.099

LIKELIHOOD WEIGHTING better than rejection sampling? \rightarrow Yes, better because uses all samples. Weakness: \rightarrow Performance bad with large number of evidence variables \rightarrow **Most samples very small weights as the number of evidence variables increases** \rightarrow The sampling process pays attention to evidence only in ancestor variables. Highly improbable event because \rightarrow Weights are low \rightarrow Problem **LL weighting works well when evidence upstream**

11. BN STRUCTURE LEARNING. Entropy: Measures the uncertainty we have about the value of a random variable. **BE CAREFUL CONDITIONAL ENTROPY $P(X, Y)$ for cond entropy is from prob of full joint table**

ENTROPY ALWAYS POSITIVE BECAUSE WE NEGATIVE LOG

- Entropy $(X) = - \sum_x p(X) \log_2 p(X)$**
- $= -(p \log p + (1-p) \log(1-p))$

Boolean $p(x1) = p$
 $p(x2) = 1-p$

ENT(X|Y) = $-\sum_y p(Y=y) ENT(X|Y=y) = -\sum_{x,y} p(X,Y) \log_2 p(X|Y)$

Structure Learning two components:

- Objective score func: score higher for better fit data, Least negative based on LogLLH
- Optimization or Search: \rightarrow Search structure with the highest score

Bayes nets are directed acyclic graphs. \rightarrow Steps for calculating Log-likelihood score

- Get the CPTs of each node using data (MLE estimates)
- Calculate the entropy of each node. 3. Use the entropies get the log-likelihood score $p(x,y)$ is from the entire joint data probability for CONDITIONAL entropy.

By adding more parents to a variable, decreases the uncertainty (entropy) \rightarrow never increase its entropy term \rightarrow We will never decrease the log-likelihood of the resulting structure. **Highest LLHOOD score is always the COMPLETE DAG \rightarrow so need penalty**

G^* is the result of adding edges to G,

- $LL(G^*|D) \geq LL(G|D)$**

AIC penalty is constant, use 1 unless BIC pen is $\log(N)/2$

Scores can be decomposed Can find BIC of a node also

ENT should be positive value. So logLL = -ENT, loglikelihood will be negative value. Scoring should be based on logLL and find highest value (least negative).

N size of DATASET

$$LL(G|Data) = -N * (ENT B + ENT C|B + ENT D|B + ENT A|C, D) = \text{--ve value}$$

Scoring measure function:

- Score(G|D) = $LL(G|D) - \varphi(N) \dim(G)$**

$\dim(G)$: dimension of BN(G)
N: Size of the dataset

First component: log-likelihood function

- Enhances the fit with data, and prefers models with more parameters

Second component: $\varphi(N) \dim(G)$

- Penalty term that favors simpler (smaller number of parameters) models
- $\dim(G)$: model complexity
- $\varphi(N)$: is a function of the data size N, weight for model complexity**

search over the space of graph structures? Move to the neighbor DAG that has the highest score \rightarrow **Neighbors are obtained through three operations by 1 step: Adding, removing, or reversing edge in one step**